

# Project 2

Dhruv Patel

16/10/2021

```
# Clean the env. variables and plots  
rm(list = ls())
```

```
library(MASS)  
library(glmnet)
```

```
## Loading required package: Matrix
```

```
## Loaded glmnet 4.1-2
```

```
library(pls)
```

```
##
```

```
## Attaching package: 'pls'
```

```
## The following object is masked from 'package:stats':
```

```
##
```

```
##      loadings
```

```
library(mgcv)
```

```
## Loading required package: nlme
```

```
## This is mgcv 1.8-36. For overview type 'help("mgcv-package")'.
```

```
library(rpart)  
library(caret)
```

```
## Loading required package: lattice
```

```
## Loading required package: ggplot2
```

```
##
```

```
## Attaching package: 'caret'
```

```
## The following object is masked from 'package:pls':
```

```
##
```

```
##      R2
```

```
library(olsrr)
```

```
##  
## Attaching package: 'olsrr'  
  
## The following object is masked from 'package:MASS':  
##  
##      cement  
  
## The following object is masked from 'package:datasets':  
##  
##      rivers
```

```
library(boot)
```

```
##  
## Attaching package: 'boot'  
  
## The following object is masked from 'package:lattice':  
##  
##      melanoma
```

```
library(dplyr)
```

```
##  
## Attaching package: 'dplyr'  
  
## The following object is masked from 'package:nlme':  
##  
##      collapse  
  
## The following object is masked from 'package:MASS':  
##  
##      select  
  
## The following objects are masked from 'package:stats':  
##  
##      filter, lag  
  
## The following objects are masked from 'package:base':  
##  
##      intersect, setdiff, setequal, union
```

```
library(e1071)  
library(ipred)  
library(performance)  
library(h2o)
```

```
##
## -----
##
## Your next step is to start H2O:
##   > h2o.init()
##
## For H2O package documentation, ask for help:
##   > ??h2o
##
## After starting H2O, you can use the Web UI at http://localhost:54321
## For more information visit https://docs.h2o.ai
##
## -----
```

```
##
## Attaching package: 'h2o'
```

```
## The following objects are masked from 'package:stats':
```

```
##
##   cor, sd, var
```

```
## The following objects are masked from 'package:base':
```

```
##
##   &&, %*%, %in%, ||, apply, as.factor, as.numeric, colnames,
##   colnames<-, ifelse, is.character, is.factor, is.numeric, log,
##   log10, log1p, log2, round, signif, trunc
```

```
rsq <- function(formula, data, indices) {
  d <- data[indices,] # allows boot to select sample
  fit <- lm(formula, data=d)
  return(summary(fit)$r.square)
}
```

```
# Helper Functions
```

```
get.MSPE = function(Y, Y.hat) {
  residuals = Y - Y.hat
  resid.sq = residuals ^ 2
  SSPE = sum(resid.sq)
  MSPE = SSPE / length(Y)
  return(MSPE)
}
```

```
data = read.csv('/Users/dhruv/Desktop/Docs/STAT_652/Project2/Data2021_final.csv', header=TRUE)
summary(data)
```

```
##           Y           X1           X2           X3
## Min.      : 7.855   Min.    :-49.00   Min.    :-48.000   Min.    :-35.40
## 1st Qu.:12.000   1st Qu.:  8.25   1st Qu.:  4.000   1st Qu.: 15.30
## Median :12.830   Median : 14.00   Median :  7.000   Median : 15.80
## Mean      :12.921   Mean     :13.61   Mean     : 7.601   Mean     :15.49
## 3rd Qu.:13.775   3rd Qu.: 19.00   3rd Qu.: 13.000   3rd Qu.: 16.30
## Max.      :17.744   Max.      :71.00   Max.      :64.000   Max.      :67.00
```

##	X4	X5	X6	X7
## Min.	:-50.000	Min. :-42.20	Min. :-41.00	Min. :-46.700
## 1st Qu.:	1.000	1st Qu.: 10.80	1st Qu.: 18.00	1st Qu.: 3.600
## Median :	2.000	Median : 12.50	Median : 21.00	Median : 3.800
## Mean :	2.348	Mean : 12.39	Mean : 21.02	Mean : 3.757
## 3rd Qu.:	3.000	3rd Qu.: 14.20	3rd Qu.: 24.00	3rd Qu.: 3.900
## Max. :	57.000	Max. : 66.10	Max. : 76.00	Max. : 54.000

##	X8	X9	X10	X11
## Min.	:-25.40	Min. :-50.00	Min. :-48.000	Min. : 159.0
## 1st Qu.:	30.30	1st Qu.: 0.00	1st Qu.: 2.000	1st Qu.: 512.0
## Median :	35.40	Median : 1.00	Median : 3.000	Median : 669.0
## Mean :	35.95	Mean : 10.39	Mean : 3.012	Mean : 711.0
## 3rd Qu.:	41.40	3rd Qu.: 5.00	3rd Qu.: 3.000	3rd Qu.: 888.8
## Max. :	97.50	Max. :149.00	Max. : 54.000	Max. :1802.0

##	X12	X13	X14	X15
## Min.	:-44.000	Min. :-49.8600	Min. : -23.20	Min. :-49.000
## 1st Qu.:	7.000	1st Qu.: 0.8825	1st Qu.: 83.67	1st Qu.: 5.000
## Median :	9.000	Median : 2.0000	Median : 189.88	Median : 9.000
## Mean :	9.311	Mean : 3.0388	Mean : 252.98	Mean : 8.749
## 3rd Qu.:	10.000	3rd Qu.: 3.9375	3rd Qu.: 346.64	3rd Qu.: 12.000
## Max. :	61.000	Max. : 58.3900	Max. :1545.88	Max. : 65.000

*# Randomly selecting 10 rows as test data from training data.*

```
set.seed(301471961)
```

```
ind = sample(nrow(data), 10, replace = TRUE)
```

```
model_test_data = data[ind, ]
```

```
data = data[-ind, ]
```

```
test_data = read.csv('/Users/dhruv/Desktop/Docs/STAT_652/Project2/Data2021test_final_noY.csv', header=TRUE)
summary(test_data)
```

##	X1	X2	X3	X4
## Min.	:-48.00	Min. :-49.000	Min. :-35.50	Min. :-50.000
## 1st Qu.:	11.00	1st Qu.: 4.000	1st Qu.: 15.30	1st Qu.: 1.000
## Median :	16.00	Median : 7.000	Median : 15.80	Median : 2.000
## Mean :	14.23	Mean : 8.007	Mean : 15.98	Mean : 1.937
## 3rd Qu.:	21.00	3rd Qu.: 13.000	3rd Qu.: 16.30	3rd Qu.: 3.000
## Max. :	74.00	Max. : 64.000	Max. : 68.50	Max. : 57.000

##	X5	X6	X7	X8
## Min.	:-43.00	Min. :-40.00	Min. :-46.800	Min. :-41.40
## 1st Qu.:	10.70	1st Qu.: 19.00	1st Qu.: 3.600	1st Qu.: 29.98
## Median :	12.40	Median : 22.00	Median : 3.800	Median : 35.80
## Mean :	12.30	Mean : 21.45	Mean : 3.631	Mean : 35.88
## 3rd Qu.:	13.93	3rd Qu.: 24.00	3rd Qu.: 3.900	3rd Qu.: 41.40
## Max. :	68.60	Max. : 79.00	Max. : 54.200	Max. :101.00

##	X9	X10	X11	X12
## Min.	:-50.000	Min. :-48.000	Min. : 23.0	Min. :-44.000
## 1st Qu.:	0.000	1st Qu.: 2.000	1st Qu.: 515.0	1st Qu.: 7.000
## Median :	1.000	Median : 3.000	Median : 695.0	Median : 9.000
## Mean :	9.049	Mean : 2.902	Mean : 719.0	Mean : 8.822
## 3rd Qu.:	4.000	3rd Qu.: 3.000	3rd Qu.: 891.2	3rd Qu.: 10.000
## Max. :	155.000	Max. : 54.000	Max. :2334.0	Max. : 61.000

##	X13	X14	X15
## Min.	:-49.9600	Min. : -40.08	Min. :-49.000

```
## 1st Qu.: 0.8275 1st Qu.: 75.11 1st Qu.: 5.000
## Median : 1.9400 Median : 168.76 Median : 9.000
## Mean : 2.4918 Mean : 236.08 Mean : 8.669
## 3rd Qu.: 3.6700 3rd Qu.: 331.84 3rd Qu.: 12.000
## Max. : 61.2000 Max. : 1954.01 Max. : 65.000
```

```
# Finding if any column has null values
sapply(data, function(x) sum(is.na(x)))
```

```
## Y X1 X2 X3 X4 X5 X6 X7 X8 X9 X10 X11 X12 X13 X14 X15
## 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
```

```
# Linear Regression Model
```

```
lm_model = lm(Y~.,data)
```

```
# For the summary we can find the X1 is the most important variable.
summary(lm_model)
```

```
##
## Call:
## lm(formula = Y ~ ., data = data)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -4.8772 -0.9091 -0.0877  0.8404  4.7129
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) 12.7248945  0.2680924  47.465 < 2e-16 ***
## X1           0.0180254  0.0040810   4.417 1.15e-05 ***
## X2          -0.0085127  0.0043564  -1.954  0.0511 .
## X3           0.0017829  0.0046490   0.384  0.7015
## X4          -0.0046188  0.0045707  -1.011  0.3126
## X5          -0.0010330  0.0045180  -0.229  0.8192
## X6          -0.0087683  0.0044878  -1.954  0.0511 .
## X7          -0.0054892  0.0047281  -1.161  0.2460
## X8           0.0017366  0.0039744   0.437  0.6623
## X9          -0.0032074  0.0021119  -1.519  0.1293
## X10          0.0075678  0.0046479   1.628  0.1039
## X11          0.0002991  0.0001903   1.571  0.1166
## X12         -0.0044165  0.0046605  -0.948  0.3436
## X13          0.0070873  0.0045852   1.546  0.1226
## X14         -0.0003014  0.0002442  -1.234  0.2175
## X15          0.0058114  0.0043543   1.335  0.1824
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 1.4 on 724 degrees of freedom
## Multiple R-squared:  0.05642,    Adjusted R-squared:  0.03687
## F-statistic: 2.886 on 15 and 724 DF,  p-value: 0.0001935
```

```
### Trying different variable selection methods to choose the important variables
```

```
# 1. Forward Selection
```

```
# Using P value
```

```
fwd_sel_model_p = ols_step_forward_p(lm_model, penter = 0.05)
```

```
# Selection Summary: X1 most important
```

```
fwd_sel_model_p
```

```
##
##                               Selection Summary
## -----
```

## Step	Variable Entered	R-Square	Adj. R-Square	C(p)	AIC	RMSE
## 1	X1	0.0250	0.0237	12.0856	2611.8281	1.4093

```
## -----
```

```
# 2. Forward Regression
```

```
# Using AIC:
```

```
fwd_sel_model_aic = ols_step_forward_aic(lm_model)
```

```
# Selection Summary: X1, X11, X9, X6, X2 are important variables
```

```
fwd_sel_model_aic
```

```
##
##                               Selection Summary
## -----
```


## Variable	AIC	Sum Sq	RSS	R-Sq	Adj. R-Sq
## X1	2611.828	37.634	1465.809	0.02503	0.02371
## X2	2610.297	44.613	1458.830	0.02967	0.02704
## X6	2608.836	51.420	1452.023	0.03420	0.03026
## X9	2607.567	57.820	1445.623	0.03846	0.03323
## X11	2606.525	63.749	1439.694	0.04240	0.03588
## X10	2606.189	68.288	1435.156	0.04542	0.03761

```
## -----
```

```
# Plotting how much each variable contributes: We can see relative contribution amongst the selected va
```

```
plot(fwd_sel_model_aic)
```

## Stepwise AIC Forward Selection



*# 3. #Backward Elimination Method*

*# Using p value*

`back_sel_model_p = ols_step_backward_p(lm_model, prem = 0.05)`

*# Elimination Summary: ["X3" "X5" "X8" "X12" "X4" "X10" "X14" "X7" "X13" "X15" "X2" "X6" "X9" "X1"]*

`back_sel_model_p`

##

##

## Elimination Summary

##	##	##	##	##	##	##	##
##	Step	Variable Removed	R-Square	Adj. R-Square	C(p)	AIC	RMSE
##	1	X5	0.0564	0.0381	14.0523	2613.6643	1.3989
##	2	X3	0.0562	0.0393	12.2055	2611.8210	1.3981
##	3	X8	0.0559	0.0403	10.3857	2610.0050	1.3973
##	4	X4	0.0548	0.0405	9.2505	2608.8879	1.3971
##	5	X12	0.0535	0.0406	8.2040	2607.8602	1.3971
##	6	X7	0.0519	0.0402	7.4950	2607.1745	1.3974
##	7	X14	0.0495	0.0391	7.2738	2606.9818	1.3981
##	8	X13	0.0476	0.0385	6.7462	2606.4743	1.3986
##	9	X15	0.0454	0.0376	6.4414	2606.1890	1.3993
##	10	X10	0.0424	0.0359	6.7576	2606.5254	1.4005
##	11	X11	0.0385	0.0332	7.7836	2607.5667	1.4024
##	12	X9	0.0342	0.0303	9.0499	2608.8356	1.4046

```
##      13      X6          0.0297      0.027      10.5241      2610.2967      1.4069
##      14      X2          0.025       0.0237      12.0856      2611.8281      1.4093
## -----
```

```
# Backward Elimination Method
```

```
# Using AIC value
```

```
back_sel_model_aic = ols_step_backward_aic(lm_model)
```

```
# Elimination Summary: ["X3" "X5" "X8" "X12" "X4" "X10" "X14" "X7" "X13" "X15"] can be removed.
```

```
back_sel_model_aic
```

```
##
```

```
##
```

```
##                               Backward Elimination Summary
```

```
## -----
## Variable      AIC      RSS      Sum Sq      R-Sq      Adj. R-Sq
## -----
## Full Model    2615.611    1418.615    84.828    0.05642    0.03687
## X5            2613.664    1418.718    84.725    0.05635    0.03813
## X3            2611.821    1419.018    84.425    0.05615    0.03925
## X8            2610.005    1419.371    84.072    0.05592    0.04034
## X4            2608.888    1421.066    82.378    0.05479    0.04051
## X12           2607.860    1422.934    80.509    0.05355    0.04057
## X7            2607.175    1425.463    77.980    0.05187    0.04018
## X14           2606.982    1428.949    74.494    0.04955    0.03915
## X13           2606.474    1431.834    71.609    0.04763    0.03852
## X15           2606.189    1435.156    68.288    0.04542    0.03761
## -----
```

```
# Step-wise using P value
```

```
stepwise_model_p = ols_step_both_p(lm_model,prem = 0.05, pent = 0.05)
```

```
# Step-wise
```

```
# Selection Summary: Only variable X1 was added
```

```
stepwise_model_p
```

```
##
```

```
##
```

```
##                               Stepwise Selection Summary
```

```
## -----
## Step  Variable      Added/      R-Square      Adj.      C(p)      AIC      RMSE
##      Removed      R-Square      R-Square
## -----
##      1      X1      addition      0.025      0.024      12.0860      2611.8281      1.4093
## -----
```

```
# Step-wise selection using AIC
```

```
stepwise_model_aic = ols_step_both_aic(lm_model)
```

```
# Step-wise Summary: ["X1" "X11" "X9" "X6" "X2"] were added to get best AIC value
```

```
stepwise_model_aic
```

```
##
```



```
##
##                               Stepwise Summary
## -----
## Variable      Method      AIC      RSS      Sum Sq      R-Sq      Adj. R-Sq
## -----
## X1            addition    2611.828    1465.809    37.634    0.02503    0.02371
## X2            addition    2610.297    1458.830    44.613    0.02967    0.02704
## X6            addition    2608.836    1452.023    51.420    0.03420    0.03026
## X9            addition    2607.567    1445.623    57.820    0.03846    0.03323
## X11           addition    2606.525    1439.694    63.749    0.04240    0.03588
## X10           addition    2606.189    1435.156    68.288    0.04542    0.03761
## -----
```

```
set.seed(301471961)
```

```
# Fitting best model using K-fold CV method
```

```
n = nrow(data)
```

```
K = 10
```

```
all.models = c("LS", "Step", "Ridge", "LAS-Min", "LAS-lse", "PLS", "GAM", "Full-Tree", "Min-Tree", "lSE-")
```

```
CV.MSPEs = array(0, dim = c(length(all.models), K))
```

```
rownames(CV.MSPEs) = all.models
```

```
colnames(CV.MSPEs) = 1:K
```

```
lambda.vals = seq(from = 0, to = 100, by = 0.05)
```

```
n = nrow(data)
```

```
for(i in 1:K){
```

```
  # Random Index
```

```
  new.order = sample.int(n)
```

```
  ind.train = which(new.order <= n * 0.75)
```

```
  ind.valid = which(new.order > n * 0.75)
```

```
  # Splitting the Data-set
```

```
  data.train = data[ind.train, ]
```

```
  data.valid = data[ind.valid, ]
```

```
  Y.train = data.train$Y
```

```
  Y.valid = data.valid$Y
```

```
  mat.train.int = model.matrix(Y ~ ., data = data.train)
```

```
  mat.train = mat.train.int[,-1]
```

```
  mat.valid.int = model.matrix(Y ~ ., data = data.valid)
```

```
  mat.valid = mat.valid.int[,-1]
```

```
  #####
```

```
  ### LS ###
```

```
  #####
```

```
  fit.ls = lm(Y ~ ., data = data.train)
```

```
  pred.ls = predict(fit.ls, data.valid)
```

```
  MSPE.ls = get.MSPE(Y.valid, pred.ls)
```

```
  CV.MSPEs["LS", i] = MSPE.ls
```

```

#####
### Step ###
#####
fit.start = lm(Y ~ 1, data = data.train)
fit.step = step(fit.start, list(upper = fit.ls), trace = 0)
pred.step = predict(fit.step, data.valid)
MSPE.step = get.MSPE(Y.valid, pred.step)
CV.MSPEs["Step", i] = MSPE.step

#####
### Ridge ###
#####

### Fit ridge regression
### We already definite lambda.vals. No need to re-invent the wheel
fit.ridge = lm.ridge(Y ~ ., lambda = lambda.vals, data = data.train)

### Get optimal lambda value
ind.min.GCV = which.min(fit.ridge$GCV)
lambda.min = lambda.vals[ind.min.GCV]

### Get coefficients for optimal model
all.coefs.ridge = coef(fit.ridge)
coef.min.ridge = all.coefs.ridge[ind.min.GCV,]

### Get predictions and MSPE on validation set
pred.ridge = mat.valid.int %% coef.min.ridge
pred.ridge = as.numeric(pred.ridge)
MSPE.ridge = get.MSPE(Y.valid, pred.ridge)
CV.MSPEs["Ridge", i] = MSPE.ridge

#####
### LASSO ###
#####

### Fit model
fit.LASSO = cv.glmnet(mat.train, Y.train)

### Get optimal lambda values
lambda.min = fit.LASSO$lambda.min
lambda.1se = fit.LASSO$lambda.1se

### Get predictions
pred.min_lasso = predict(fit.LASSO, mat.valid, lambda.min)
pred.1se = predict(fit.LASSO, mat.valid, lambda.1se)

### Get and store MSPEs
MSPE.min = get.MSPE(Y.valid, pred.min_lasso)
MSPE.1se = get.MSPE(Y.valid, pred.1se)
CV.MSPEs["LAS-Min", i] = MSPE.min
CV.MSPEs["LAS-1se", i] = MSPE.1se

```

```
#####
### Partial Least Squares ###
#####

### Fit PLS
fit.pls = plsr(Y ~ ., data = data.train, validation = "CV", segments = 10)

### Get optimal number of folds
CV.pls = fit.pls$validation
PRESS.pls = CV.pls$PRESS
n.comps = which.min(PRESS.pls)

### Get predictions and MSPE
pred.pls = predict(fit.pls, data.valid, ncomp = n.comps)
MSPE.pls = get.MSPE(Y.valid, pred.pls)
CV.MSPEs["PLS", i] = MSPE.pls

#####
### GAM ###
#####

### Fit model
fit.gam = gam(Y ~ s(X1) + s(X11) + s(X9) + s(X6) + s(X2) + s(X10) + s(X15), data = data.train)

### Get predictions and MSPE
pred.gam = predict(fit.gam, data.valid)
MSPE.gam = get.MSPE(Y.valid, pred.gam)
CV.MSPEs["GAM", i] = MSPE.gam

#####
### Full Tree ###
#####

fit.tree = rpart(Y ~ ., data = data.train, cp = 0)

### Get the CP table
info.tree = fit.tree$cptable

### Get predictions
pred.full = predict(fit.tree, data.valid)
MSPE.full = get.MSPE(Y.valid, pred.full)
CV.MSPEs["Full-Tree", i] = MSPE.full

#####
### Min CV Tree ###
#####

### Get minimum CV error and corresponding CP value
ind.best = which.min(info.tree[, "xerror"])
CV.best = info.tree[ind.best, "xerror"]
CP.best = info.tree[ind.best, "CP"]
```

```

### Get the geometric mean of best CP with one above it
if (ind.best == 1) {
  CP.GM = CP.best
}
else{
  CP.above = info.tree[ind.best - 1, "CP"]
  CP.GM = sqrt(CP.best * CP.above)
}

### Fit minimum CV error tree
fit.tree.min = prune(fit.tree, cp = CP.best)

### Get predictions and MSPE
pred.min = predict(fit.tree.min, data.valid)
MSPE.min = get.MSPE(Y.valid, pred.min)
CV.MSPEs["Min-Tree", i] = MSPE.min

#####
### 1SE Rule CV Tree ###
#####

### Get 1se rule CP value
err.min = info.tree[ind.best, "xerror"]
se.min = info.tree[ind.best, "xstd"]

threshold = err.min + se.min
ind.1se = min(which(info.tree[1:ind.best, "xerror"] < threshold))

### Take geometric mean with superior row
CP.1se.raw = info.tree[ind.1se, "CP"]

if (ind.1se == 1) {
  ### If best CP is in row 1, store this value
  CP.1se = CP.1se.raw
}
else{
  ### If best CP is not in row 1, average this with the value from the ### row above it.
  ### Value from row above
  CP.above = info.tree[ind.1se - 1, "CP"]

  ### (Geometric) average
  CP.1se = sqrt(CP.1se.raw * CP.above)
}

### Prune the tree
fit.tree.1se = prune(fit.tree, cp = CP.1se)

### Get predictions and MSPE
pred.1se = predict(fit.tree.1se, data.valid)
MSPE.1se = get.MSPE(Y.valid, pred.1se)
CV.MSPEs["1SE-Tree", i] = MSPE.1se
}

```

## CV.MSPEs

```
##           1           2           3           4           5           6           7
## LS       2.262826 2.199506 2.174115 2.239086 2.026980 1.644678 2.041498
## Step     2.287486 2.239805 2.189244 2.258342 2.064131 1.699231 2.065840
## Ridge    2.254373 2.203214 2.151474 2.227014 1.997511 1.628392 2.027461
## LAS-Min  2.254025 2.262128 2.144635 2.245310 2.016814 1.641607 2.048312
## LAS-1se  2.309883 2.316721 2.172645 2.287196 1.989073 1.652671 2.074964
## PLS      2.331110 2.285897 2.173540 2.217923 2.038460 1.647188 2.040617
## GAM      2.091092 1.933269 1.882283 2.059659 2.022764 1.631334 2.022538
## Full-Tree 2.868454 2.343721 2.282801 2.712111 2.513663 2.474951 2.741155
## Min-Tree 2.304910 2.393777 2.036022 2.235144 2.075740 1.789475 2.082046
## 1SE-Tree 2.309883 2.130736 1.932816 2.287196 2.056348 1.792475 2.074964
##           8           9          10
## LS       1.892345 2.086754 2.002616
## Step     1.914008 2.063151 2.038700
## Ridge    1.886477 2.069589 1.990547
## LAS-Min  1.892848 2.068131 2.013093
## LAS-1se  1.931517 2.099387 2.044572
## PLS      1.884597 2.060205 2.060292
## GAM      1.694160 1.758533 1.908523
## Full-Tree 1.909727 2.558427 2.378676
## Min-Tree 1.824420 1.938802 2.028302
## 1SE-Tree 1.931517 2.099387 2.044572
```

## # Average MSPE


```
rowMeans(CV.MSPEs)
```

```
##      LS      Step      Ridge  LAS-Min  LAS-1se      PLS      GAM Full-Tree
## 2.057040 2.081994 2.043605 2.058690 2.087863 2.073983 1.900415 2.478369
## Min-Tree 1SE-Tree
## 2.070864 2.065989
```

## ### Make boxplot


```
boxplot(CV.MSPEs, las = 2, main = "MSPE Boxplot")
```

## MSPE Boxplot



```
### Get relative MSPEs and make boxplot
CV.RMSPEs = apply(CV.MSPEs, 1, function(W) W/min(W))
CV.RMSPEs = t(CV.RMSPEs)
boxplot(CV.RMSPEs, las = 2, main = "RMSPE Boxplot")
```

## RMSPE Boxplot



```
# Code for best model
```

```
# Initializing environment
```

```
h2o.init(nthreads = -1)
```

```
## Connection successful!
```

```
##
```

```
## R is connected to the H2O cluster:
```

```
## H2O cluster uptime: 2 hours 4 minutes
```

```
## H2O cluster timezone: America/Vancouver
```

```
## H2O data parsing timezone: UTC
```

```
## H2O cluster version: 3.34.0.3
```

```
## H2O cluster version age: 1 month and 21 days
```

```
## H2O cluster name: H2O_started_from_R_dhruv_ygi863
```

```
## H2O cluster total nodes: 1
```

```
## H2O cluster total memory: 0.70 GB
```

```
## H2O cluster total cores: 4
```

```
## H2O cluster allowed cores: 4
```

```
## H2O cluster healthy: TRUE
```

```
## H2O Connection ip: localhost
```

```
## H2O Connection port: 54321
```

```
## H2O Connection proxy: NA
```

```
## H2O Internal Security: FALSE
```

```
## H2O API Extensions: Amazon S3, XGBoost, Algos, AutoML, Core V3, TargetEncoder, Core V4
```

```
## R Version: R version 4.1.1 (2021-08-10)
```

```

y = "y"
x = setdiff(names(data), y)

train.h2o = as.h2o(data)

```

```
## |
```

```

h2o.fit1 <- h2o.gbm(
  x = x,
  y = y,
  training_frame = train.h2o,
  nfolds = 5,
  ntrees = 5000,
  stopping_rounds = 10,
  stopping_tolerance = 0,
  seed = 301471961
)

```

```
## Warning in .h2o.processResponseWarnings(res): early stopping is enabled but neither score_tree_inter
```

```
## |
```

```
h2o.fit1@parameters$ntrees # 41
```

```
## [1] 39
```

```
h2o.rmse(h2o.fit1, xval = TRUE) # 1.278083
```

```
## [1] 1.278757
```

```

split = h2o.splitFrame(train.h2o, ratios = 0.75)
train = split[[1]]
valid = split[[2]]

```

```

h2o.final <- h2o.gbm(
  x = x,
  y = y,
  training_frame = train.h2o,
  nfolds = 5,
  ntrees = 5000,
  learn_rate = 0.01,
  learn_rate_annealing = 1,
  max_depth = 1,
  min_rows = 1,
  sample_rate = 0.75,
  col_sample_rate = 1,
  stopping_rounds = 10,
  stopping_tolerance = 0,
  seed = 301471961
)

```



```
## Warning in .h2o.processResponseWarnings(res): early stopping is enabled but neither score_tree_inter
```

```
##      |
```

```
h2o.final@parameters$ntrees
```


```
## [1] 5000
```

```
h2o.rmse(h2o.final, xval = TRUE)
```

```
## [1] 1.294353
```

```
h2o.varimp_plot(h2o.final, num_of_features = 10)
```

### Variable Importance: GBM



```
# Validating model
```

```
test.h2o <- as.h2o(model_test_data)
```

```
##      |
```

```
h2o.performance(model = h2o.final, newdata = test.h2o)
```

```
## H2ORegressionMetrics: gbm
```

```
##
```

```
## MSE: 3.20338
```

```
## RMSE: 1.789799
```

```
## MAE: 1.506297
```

```
## RMSLE: 0.1345529
```

```
## Mean Residual Deviance : 3.20338
```

```
h2o.predict(h2o.final, newdata = test.h2o)
```

```
## |
```

```
## predict
```

```
## 1 12.71176
```

```
## 2 10.61467
```

```
## 3 12.84556
```

```
## 4 13.57291
```

```
## 5 13.08423
```

```
## 6 12.92769
```

```
##
```

```
## [10 rows x 1 column]
```

```
pred.h2o = predict(h2o.final, test.h2o)
```

```
## |
```

```
test_acc = get.MSPE(model_test_data$Y, pred.h2o) # 1.10928
```

```
# Getting prediction for test_data
```

```
test.h2o <- as.h2o(test_data)
```

```
## |
```

```
h2o.performance(model = h2o.final, newdata = test.h2o)
```

```
## [1] "WARNING: Model metrics cannot be calculated and metric_json is empty due to the absence of the "
```

```
## NULL
```

```
h2o.predict(h2o.final, newdata = test.h2o)
```

```
## |
```

```
## predict
```

```
## 1 12.11686
```

```
## 2 13.00003
```

```
## 3 13.12212
```

```
## 4 13.75701
```

```
## 5 12.61131
```

```
## 6 12.81275
```

```
##
```

```
## [3000 rows x 1 column]
```

```
predictions = predict(h2o.final, test.h2o)
```

```
## |
```

```
h2o.exportFile(predictions, path = "test.csv")
```

```
## |
```