

STAT 652 - Project 2

Dhruv Patel, 301471961

1. What models or machines did you attempt to fit?

Answer:

- **LS**

```
fit.ls = lm(Y ~ ., data = data.train)
```

- **Step**

```
fit.step = step(fit.start, list(upper = fit.ls), trace = 0)
```

- **Ridge**

```
fit.ridge = lm.ridge(Y ~ ., lambda = lambda.vals, data = data.train)
```

- **LASSO**

```
fit.LASSO = cv.glmnet(mat.train, Y.train)
```

- **LAS-Min**

```
lambda.min = fit.LASSO$lambda.min
```

- **LAS-1se**

```
lambda.1se = fit.LASSO$lambda.1se
```

- **PLS**

```
fit.pls = pls(Y ~ ., data = data.train, validation = "CV", segments = 10)
```

- **GAM**

```
fit.gam = gam(Y ~ s(X1) + s(X11) + s(X9) + s(X6) + s(X2) + s(X10) + s(X15), data =  
data.train)
```

- **Full-Tree**

```
fit.tree.min = prune(fit.tree, cp = CP.best)
```

- **Min-Tree**

```
fit.tree.1se = prune(fit.tree, cp = CP.1se)
```

- **1SE-Tree**

```
fit.tree = rpart(Y ~ ., data = data.train, cp = 0)
```

- **h2o-GBM**

```
h2o.final <- h2o.gbm( x = x, y = y, training_frame = train.h2o,  
nfold = 5, ntrees = 5000, learn_rate = 0.01,  
learn_rate_annealing = 1, max_depth = 1, min_rows = 1, sample_rate  
= 0.75, col_sample_rate = 1, stopping_rounds = 10,  
stopping_tolerance = 0, seed = 301471961 )
```

2. What process(es) did you use to evaluate and compare models and to select your final model?

Answer: I started by cleaning and imputing values since we had less training data. Also, removing the outliers to model doesn't train on them. Further, I randomly selected 10 rows for test data. And trained the model on 740 rows, which is split into train data (75%) and validation data (25%). Then, I ran several selection variable techniques to find important variables, this will not only reduce the dimensionality but also prevent the model from over fitting. Then applied a 10-fold cross validation (evaluation method) on the shuffled dataset and MSPE (mean square prediction error) to compare different models. And based on the MSPE the model with minimum MSPE was selected.

3. Did you tune any methods?

Answer: Yes, I ran a full cartesian grid search that examined each and every combination for the best tuning for the model. Hyper-parameters such as max_depth, min_rows, learn_rate, sample_rate etc. In the best model [H2o-GBM] which had an RMSE = 1.294353


4. What was your chosen prediction machine?

Answer: I chose [H2o-GBM] as the best model with seed 301471961. Because, this model performed well compared to others on validation data and even the error for the test data was low. Which points out that the model did not overfit. Hence, the prediction was done using GAM on the important selected variables.

```
h2o.final <- h2o.gbm(  
  x = x,  
  y = y,  
  training_frame = train.h2o,  
  nfolds = 5,  
  ntrees = 5000,  
  learn_rate = 0.01,  
  learn_rate_annealing = 1,  
  max_depth = 1,  
  min_rows = 1,  
  sample_rate = 0.75,  
  col_sample_rate = 1,  
  stopping_rounds = 10,  
  stopping_tolerance = 0,  
  seed = 301471961  
)
```

5. List the variables that you believe are important?

Answer: Based on the selected best model top 10 important variables are as follows
X15, X1, X12, X2, X4, X10, X6, X13, X8, X11.



Note: Attaching the pdf of the rmd file for the code.

Project 2

Dhruv Patel

16/10/2021

```
# Clean the env. variables and plots  
rm(list = ls())
```

```
library(MASS)  
library(glmnet)
```

```
## Loading required package: Matrix
```

```
## Loaded glmnet 4.1-2
```

```
library(pls)
```

```
##
```

```
## Attaching package: 'pls'
```

```
## The following object is masked from 'package:stats':
```

```
##
```

```
##      loadings
```

```
library(mgcv)
```

```
## Loading required package: nlme
```

```
## This is mgcv 1.8-36. For overview type 'help("mgcv-package")'.
```

```
library(rpart)  
library(caret)
```

```
## Loading required package: lattice
```

```
## Loading required package: ggplot2
```

```
##
```

```
## Attaching package: 'caret'
```

```
## The following object is masked from 'package:pls':
```

```
##
```

```
##      R2
```

```
library(olsrr)
```

```
##  
## Attaching package: 'olsrr'  
  
## The following object is masked from 'package:MASS':  
##  
##      cement  
  
## The following object is masked from 'package:datasets':  
##  
##      rivers
```

```
library(boot)
```

```
##  
## Attaching package: 'boot'  
  
## The following object is masked from 'package:lattice':  
##  
##      melanoma
```

```
library(dplyr)
```

```
##  
## Attaching package: 'dplyr'  
  
## The following object is masked from 'package:nlme':  
##  
##      collapse  
  
## The following object is masked from 'package:MASS':  
##  
##      select  
  
## The following objects are masked from 'package:stats':  
##  
##      filter, lag  
  
## The following objects are masked from 'package:base':  
##  
##      intersect, setdiff, setequal, union
```

```
library(e1071)  
library(ipred)  
library(performance)  
library(h2o)
```

```
##
## -----
##
## Your next step is to start H2O:
##   > h2o.init()
##
## For H2O package documentation, ask for help:
##   > ??h2o
##
## After starting H2O, you can use the Web UI at http://localhost:54321
## For more information visit https://docs.h2o.ai
##
## -----
```

```
##
## Attaching package: 'h2o'
```

```
## The following objects are masked from 'package:stats':
```

```
##
##   cor, sd, var
```

```
## The following objects are masked from 'package:base':
```

```
##
##   &&, %*%, %in%, ||, apply, as.factor, as.numeric, colnames,
##   colnames<-, ifelse, is.character, is.factor, is.numeric, log,
##   log10, log1p, log2, round, signif, trunc
```

```
rsq <- function(formula, data, indices) {
  d <- data[indices,] # allows boot to select sample
  fit <- lm(formula, data=d)
  return(summary(fit)$r.square)
}
```

```
# Helper Functions
```

```
get.MSPE = function(Y, Y.hat) {
  residuals = Y - Y.hat
  resid.sq = residuals ^ 2
  SSPE = sum(resid.sq)
  MSPE = SSPE / length(Y)
  return(MSPE)
}
```

```
data = read.csv('/Users/dhruv/Desktop/Docs/STAT_652/Project2/Data2021_final.csv', header=TRUE)
summary(data)
```

```
##           Y           X1           X2           X3
## Min.      : 7.855   Min.    :-49.00   Min.    :-48.000   Min.    :-35.40
## 1st Qu.:12.000   1st Qu.:  8.25   1st Qu.:  4.000   1st Qu.: 15.30
## Median :12.830   Median : 14.00   Median :  7.000   Median : 15.80
## Mean      :12.921   Mean     :13.61   Mean     : 7.601   Mean     :15.49
## 3rd Qu.:13.775   3rd Qu.: 19.00   3rd Qu.: 13.000   3rd Qu.: 16.30
## Max.      :17.744   Max.      :71.00   Max.      :64.000   Max.      :67.00
```

	X4	X5	X6	X7
## Min.	:-50.000	Min. :-42.20	Min. :-41.00	Min. :-46.700
## 1st Qu.:	1.000	1st Qu.: 10.80	1st Qu.: 18.00	1st Qu.: 3.600
## Median :	2.000	Median : 12.50	Median : 21.00	Median : 3.800
## Mean :	2.348	Mean : 12.39	Mean : 21.02	Mean : 3.757
## 3rd Qu.:	3.000	3rd Qu.: 14.20	3rd Qu.: 24.00	3rd Qu.: 3.900
## Max. :	57.000	Max. : 66.10	Max. : 76.00	Max. : 54.000

	X8	X9	X10	X11
## Min.	:-25.40	Min. :-50.00	Min. :-48.000	Min. : 159.0
## 1st Qu.:	30.30	1st Qu.: 0.00	1st Qu.: 2.000	1st Qu.: 512.0
## Median :	35.40	Median : 1.00	Median : 3.000	Median : 669.0
## Mean :	35.95	Mean : 10.39	Mean : 3.012	Mean : 711.0
## 3rd Qu.:	41.40	3rd Qu.: 5.00	3rd Qu.: 3.000	3rd Qu.: 888.8
## Max. :	97.50	Max. :149.00	Max. : 54.000	Max. :1802.0

	X12	X13	X14	X15
## Min.	:-44.000	Min. :-49.8600	Min. : -23.20	Min. :-49.000
## 1st Qu.:	7.000	1st Qu.: 0.8825	1st Qu.: 83.67	1st Qu.: 5.000
## Median :	9.000	Median : 2.0000	Median : 189.88	Median : 9.000
## Mean :	9.311	Mean : 3.0388	Mean : 252.98	Mean : 8.749
## 3rd Qu.:	10.000	3rd Qu.: 3.9375	3rd Qu.: 346.64	3rd Qu.: 12.000
## Max. :	61.000	Max. : 58.3900	Max. :1545.88	Max. : 65.000

Randomly selecting 10 rows as test data from training data.

```
set.seed(301471961)
```

```
ind = sample(nrow(data), 10, replace = TRUE)
```

```
model_test_data = data[ind, ]
```

```
data = data[-ind, ]
```

```
test_data = read.csv('/Users/dhruv/Desktop/Docs/STAT_652/Project2/Data2021test_final_noY.csv', header=TRUE)
summary(test_data)
```

	X1	X2	X3	X4
## Min.	:-48.00	Min. :-49.000	Min. :-35.50	Min. :-50.000
## 1st Qu.:	11.00	1st Qu.: 4.000	1st Qu.: 15.30	1st Qu.: 1.000
## Median :	16.00	Median : 7.000	Median : 15.80	Median : 2.000
## Mean :	14.23	Mean : 8.007	Mean : 15.98	Mean : 1.937
## 3rd Qu.:	21.00	3rd Qu.: 13.000	3rd Qu.: 16.30	3rd Qu.: 3.000
## Max. :	74.00	Max. : 64.000	Max. : 68.50	Max. : 57.000

	X5	X6	X7	X8
## Min.	:-43.00	Min. :-40.00	Min. :-46.800	Min. :-41.40
## 1st Qu.:	10.70	1st Qu.: 19.00	1st Qu.: 3.600	1st Qu.: 29.98
## Median :	12.40	Median : 22.00	Median : 3.800	Median : 35.80
## Mean :	12.30	Mean : 21.45	Mean : 3.631	Mean : 35.88
## 3rd Qu.:	13.93	3rd Qu.: 24.00	3rd Qu.: 3.900	3rd Qu.: 41.40
## Max. :	68.60	Max. : 79.00	Max. : 54.200	Max. :101.00

	X9	X10	X11	X12
## Min.	:-50.000	Min. :-48.000	Min. : 23.0	Min. :-44.000
## 1st Qu.:	0.000	1st Qu.: 2.000	1st Qu.: 515.0	1st Qu.: 7.000
## Median :	1.000	Median : 3.000	Median : 695.0	Median : 9.000
## Mean :	9.049	Mean : 2.902	Mean : 719.0	Mean : 8.822
## 3rd Qu.:	4.000	3rd Qu.: 3.000	3rd Qu.: 891.2	3rd Qu.: 10.000
## Max. :	155.000	Max. : 54.000	Max. :2334.0	Max. : 61.000

	X13	X14	X15
## Min.	:-49.9600	Min. : -40.08	Min. :-49.000

```
## 1st Qu.: 0.8275 1st Qu.: 75.11 1st Qu.: 5.000
## Median : 1.9400 Median : 168.76 Median : 9.000
## Mean : 2.4918 Mean : 236.08 Mean : 8.669
## 3rd Qu.: 3.6700 3rd Qu.: 331.84 3rd Qu.: 12.000
## Max. : 61.2000 Max. : 1954.01 Max. : 65.000
```

```
# Finding if any column has null values
sapply(data, function(x) sum(is.na(x)))
```

```
## Y X1 X2 X3 X4 X5 X6 X7 X8 X9 X10 X11 X12 X13 X14 X15
## 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
```

```
# Linear Regression Model
```

```
lm_model = lm(Y~.,data)
```

```
# For the summary we can find the X1 is the most important variable.
summary(lm_model)
```

```
##
## Call:
## lm(formula = Y ~ ., data = data)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -4.8772 -0.9091 -0.0877  0.8404  4.7129
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) 12.7248945  0.2680924  47.465 < 2e-16 ***
## X1           0.0180254  0.0040810   4.417 1.15e-05 ***
## X2          -0.0085127  0.0043564  -1.954  0.0511 .
## X3           0.0017829  0.0046490   0.384  0.7015
## X4          -0.0046188  0.0045707  -1.011  0.3126
## X5          -0.0010330  0.0045180  -0.229  0.8192
## X6          -0.0087683  0.0044878  -1.954  0.0511 .
## X7          -0.0054892  0.0047281  -1.161  0.2460
## X8           0.0017366  0.0039744   0.437  0.6623
## X9          -0.0032074  0.0021119  -1.519  0.1293
## X10          0.0075678  0.0046479   1.628  0.1039
## X11          0.0002991  0.0001903   1.571  0.1166
## X12          -0.0044165  0.0046605  -0.948  0.3436
## X13          0.0070873  0.0045852   1.546  0.1226
## X14          -0.0003014  0.0002442  -1.234  0.2175
## X15          0.0058114  0.0043543   1.335  0.1824
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 1.4 on 724 degrees of freedom
## Multiple R-squared:  0.05642, Adjusted R-squared:  0.03687
## F-statistic: 2.886 on 15 and 724 DF, p-value: 0.0001935
```



```
### Trying different variable selection methods to choose the important variables
```

```
# 1. Forward Selection
```

```
# Using P value
```

```
fwd_sel_model_p = ols_step_forward_p(lm_model, penter = 0.05)
```

```
# Selection Summary: X1 most important
```

```
fwd_sel_model_p
```

```
##
##                               Selection Summary
## -----
```

## Step	Variable Entered	R-Square	Adj. R-Square	C(p)	AIC	RMSE
## 1	X1	0.0250	0.0237	12.0856	2611.8281	1.4093

```
## -----
```

```
# 2. Forward Regression
```

```
# Using AIC:
```

```
fwd_sel_model_aic = ols_step_forward_aic(lm_model)
```

```
# Selection Summary: X1, X11, X9, X6, X2 are important variables
```

```
fwd_sel_model_aic
```

```
##
##                               Selection Summary
## -----
```


## Variable	AIC	Sum Sq	RSS	R-Sq	Adj. R-Sq
## X1	2611.828	37.634	1465.809	0.02503	0.02371
## X2	2610.297	44.613	1458.830	0.02967	0.02704
## X6	2608.836	51.420	1452.023	0.03420	0.03026
## X9	2607.567	57.820	1445.623	0.03846	0.03323
## X11	2606.525	63.749	1439.694	0.04240	0.03588
## X10	2606.189	68.288	1435.156	0.04542	0.03761

```
## -----
```

```
# Plotting how much each variable contributes: We can see relative contribution amongst the selected va
```

```
plot(fwd_sel_model_aic)
```

Stepwise AIC Forward Selection



3. #Backward Elimination Method

Using p value

```
back_sel_model_p = ols_step_backward_p(lm_model, prem = 0.05)
```

Elimination Summary: ["X3" "X5" "X8" "X12" "X4" "X10" "X14" "X7" "X13" "X15" "X2" "X6" "X9" "X1"]

```
back_sel_model_p
```

##

##

Elimination Summary

##	##	##	##	##	##	##	##
##	Step	Variable Removed	R-Square	Adj. R-Square	C(p)	AIC	RMSE
##	1	X5	0.0564	0.0381	14.0523	2613.6643	1.3989
##	2	X3	0.0562	0.0393	12.2055	2611.8210	1.3981
##	3	X8	0.0559	0.0403	10.3857	2610.0050	1.3973
##	4	X4	0.0548	0.0405	9.2505	2608.8879	1.3971
##	5	X12	0.0535	0.0406	8.2040	2607.8602	1.3971
##	6	X7	0.0519	0.0402	7.4950	2607.1745	1.3974
##	7	X14	0.0495	0.0391	7.2738	2606.9818	1.3981
##	8	X13	0.0476	0.0385	6.7462	2606.4743	1.3986
##	9	X15	0.0454	0.0376	6.4414	2606.1890	1.3993
##	10	X10	0.0424	0.0359	6.7576	2606.5254	1.4005
##	11	X11	0.0385	0.0332	7.7836	2607.5667	1.4024
##	12	X9	0.0342	0.0303	9.0499	2608.8356	1.4046

```
##    13    X6          0.0297      0.027      10.5241      2610.2967      1.4069
##    14    X2          0.025       0.0237      12.0856      2611.8281      1.4093
## -----
```

```
# Backward Elimination Method
```

```
# Using AIC value
```

```
back_sel_model_aic = ols_step_backward_aic(lm_model)
```

```
# Elimination Summary: ["X3" "X5" "X8" "X12" "X4" "X10" "X14" "X7" "X13" "X15"] can be removed.
```

```
back_sel_model_aic
```

```
##
```

```
##
```

```
##                               Backward Elimination Summary
```

```
## -----
## Variable      AIC      RSS      Sum Sq      R-Sq      Adj. R-Sq
## -----
## Full Model    2615.611    1418.615    84.828    0.05642    0.03687
## X5            2613.664    1418.718    84.725    0.05635    0.03813
## X3            2611.821    1419.018    84.425    0.05615    0.03925
## X8            2610.005    1419.371    84.072    0.05592    0.04034
## X4            2608.888    1421.066    82.378    0.05479    0.04051
## X12           2607.860    1422.934    80.509    0.05355    0.04057
## X7            2607.175    1425.463    77.980    0.05187    0.04018
## X14           2606.982    1428.949    74.494    0.04955    0.03915
## X13           2606.474    1431.834    71.609    0.04763    0.03852
## X15           2606.189    1435.156    68.288    0.04542    0.03761
## -----
```

```
# Step-wise using P value
```

```
stepwise_model_p = ols_step_both_p(lm_model,prem = 0.05, pent = 0.05)
```

```
# Step-wise
```

```
# Selection Summary: Only variable X1 was added
```

```
stepwise_model_p
```

```
##
```

```
##                               Stepwise Selection Summary
```

```
## -----
## Step  Variable      Added/      R-Square      Adj.      C(p)      AIC      RMSE
##      Removed
## 1      X1      addition      0.025      0.024      12.0860      2611.8281      1.4093
## -----
```

```
# Step-wise selection using AIC
```

```
stepwise_model_aic = ols_step_both_aic(lm_model)
```

```
# Step-wise Summary: ["X1" "X11" "X9" "X6" "X2"] were added to get best AIC value
```

```
stepwise_model_aic
```

```
##
```

```
##
##                               Stepwise Summary
## -----
## Variable      Method      AIC      RSS      Sum Sq      R-Sq      Adj. R-Sq
## -----
## X1            addition    2611.828    1465.809    37.634    0.02503    0.02371
## X2            addition    2610.297    1458.830    44.613    0.02967    0.02704
## X6            addition    2608.836    1452.023    51.420    0.03420    0.03026
## X9            addition    2607.567    1445.623    57.820    0.03846    0.03323
## X11           addition    2606.525    1439.694    63.749    0.04240    0.03588
## X10           addition    2606.189    1435.156    68.288    0.04542    0.03761
## -----
```

```
set.seed(301471961)
```

```
# Fitting best model using K-fold CV method
```

```
n = nrow(data)
```

```
K = 10
```

```
all.models = c("LS", "Step", "Ridge", "LAS-Min", "LAS-lse", "PLS", "GAM", "Full-Tree", "Min-Tree", "lSE-")
```

```
CV.MSPEs = array(0, dim = c(length(all.models), K))
```

```
rownames(CV.MSPEs) = all.models
```

```
colnames(CV.MSPEs) = 1:K
```

```
lambda.vals = seq(from = 0, to = 100, by = 0.05)
```

```
n = nrow(data)
```

```
for(i in 1:K){
```

```
  # Random Index
```

```
  new.order = sample.int(n)
```

```
  ind.train = which(new.order <= n * 0.75)
```

```
  ind.valid = which(new.order > n * 0.75)
```

```
  # Splitting the Data-set
```

```
  data.train = data[ind.train, ]
```

```
  data.valid = data[ind.valid, ]
```

```
  Y.train = data.train$Y
```

```
  Y.valid = data.valid$Y
```

```
  mat.train.int = model.matrix(Y ~ ., data = data.train)
```

```
  mat.train = mat.train.int[,-1]
```

```
  mat.valid.int = model.matrix(Y ~ ., data = data.valid)
```

```
  mat.valid = mat.valid.int[,-1]
```

```
  #####
```

```
  ### LS ###
```

```
  #####
```

```
  fit.ls = lm(Y ~ ., data = data.train)
```

```
  pred.ls = predict(fit.ls, data.valid)
```

```
  MSPE.ls = get.MSPE(Y.valid, pred.ls)
```

```
  CV.MSPEs["LS", i] = MSPE.ls
```

```

#####
### Step ###
#####
fit.start = lm(Y ~ 1, data = data.train)
fit.step = step(fit.start, list(upper = fit.ls), trace = 0)
pred.step = predict(fit.step, data.valid)
MSPE.step = get.MSPE(Y.valid, pred.step)
CV.MSPEs["Step", i] = MSPE.step

#####
### Ridge ###
#####

### Fit ridge regression
### We already definite lambda.vals. No need to re-invent the wheel
fit.ridge = lm.ridge(Y ~ ., lambda = lambda.vals, data = data.train)

### Get optimal lambda value
ind.min.GCV = which.min(fit.ridge$GCV)
lambda.min = lambda.vals[ind.min.GCV]

### Get coefficients for optimal model
all.coefs.ridge = coef(fit.ridge)
coef.min.ridge = all.coefs.ridge[ind.min.GCV,]

### Get predictions and MSPE on validation set
pred.ridge = mat.valid.int %% coef.min.ridge
pred.ridge = as.numeric(pred.ridge)
MSPE.ridge = get.MSPE(Y.valid, pred.ridge)
CV.MSPEs["Ridge", i] = MSPE.ridge

#####
### LASSO ###
#####

### Fit model
fit.LASSO = cv.glmnet(mat.train, Y.train)

### Get optimal lambda values
lambda.min = fit.LASSO$lambda.min
lambda.1se = fit.LASSO$lambda.1se

### Get predictions
pred.min_lasso = predict(fit.LASSO, mat.valid, lambda.min)
pred.1se = predict(fit.LASSO, mat.valid, lambda.1se)

### Get and store MSPEs
MSPE.min = get.MSPE(Y.valid, pred.min_lasso)
MSPE.1se = get.MSPE(Y.valid, pred.1se)
CV.MSPEs["LAS-Min", i] = MSPE.min
CV.MSPEs["LAS-1se", i] = MSPE.1se

```

```
#####
### Partial Least Squares ###
#####

### Fit PLS
fit.pls = plsr(Y ~ ., data = data.train, validation = "CV", segments = 10)

### Get optimal number of folds
CV.pls = fit.pls$validation
PRESS.pls = CV.pls$PRESS
n.comps = which.min(PRESS.pls)

### Get predictions and MSPE
pred.pls = predict(fit.pls, data.valid, ncomp = n.comps)
MSPE.pls = get.MSPE(Y.valid, pred.pls)
CV.MSPEs["PLS", i] = MSPE.pls

#####
### GAM ###
#####

### Fit model
fit.gam = gam(Y ~ s(X1) + s(X11) + s(X9) + s(X6) + s(X2) + s(X10) + s(X15), data = data.train)

### Get predictions and MSPE
pred.gam = predict(fit.gam, data.valid)
MSPE.gam = get.MSPE(Y.valid, pred.gam)
CV.MSPEs["GAM", i] = MSPE.gam

#####
### Full Tree ###
#####

fit.tree = rpart(Y ~ ., data = data.train, cp = 0)

### Get the CP table
info.tree = fit.tree$cptable

### Get predictions
pred.full = predict(fit.tree, data.valid)
MSPE.full = get.MSPE(Y.valid, pred.full)
CV.MSPEs["Full-Tree", i] = MSPE.full

#####
### Min CV Tree ###
#####

### Get minimum CV error and corresponding CP value
ind.best = which.min(info.tree[, "xerror"])
CV.best = info.tree[ind.best, "xerror"]
CP.best = info.tree[ind.best, "CP"]
```

```

### Get the geometric mean of best CP with one above it
if (ind.best == 1) {
  CP.GM = CP.best
}
else{
  CP.above = info.tree[ind.best - 1, "CP"]
  CP.GM = sqrt(CP.best * CP.above)
}

### Fit minimum CV error tree
fit.tree.min = prune(fit.tree, cp = CP.best)

### Get predictions and MSPE
pred.min = predict(fit.tree.min, data.valid)
MSPE.min = get.MSPE(Y.valid, pred.min)
CV.MSPEs["Min-Tree", i] = MSPE.min

#####
### 1SE Rule CV Tree ###
#####

### Get 1se rule CP value
err.min = info.tree[ind.best, "xerror"]
se.min = info.tree[ind.best, "xstd"]

threshold = err.min + se.min
ind.1se = min(which(info.tree[1:ind.best, "xerror"] < threshold))

### Take geometric mean with superior row
CP.1se.raw = info.tree[ind.1se, "CP"]

if (ind.1se == 1) {
  ### If best CP is in row 1, store this value
  CP.1se = CP.1se.raw
}
else{
  ### If best CP is not in row 1, average this with the value from the ### row above it.
  ### Value from row above
  CP.above = info.tree[ind.1se - 1, "CP"]

  ### (Geometric) average
  CP.1se = sqrt(CP.1se.raw * CP.above)
}

### Prune the tree
fit.tree.1se = prune(fit.tree, cp = CP.1se)

### Get predictions and MSPE
pred.1se = predict(fit.tree.1se, data.valid)
MSPE.1se = get.MSPE(Y.valid, pred.1se)
CV.MSPEs["1SE-Tree", i] = MSPE.1se
}

```

CV.MSPEs

```
##           1           2           3           4           5           6           7
## LS       2.262826 2.199506 2.174115 2.239086 2.026980 1.644678 2.041498
## Step     2.287486 2.239805 2.189244 2.258342 2.064131 1.699231 2.065840
## Ridge    2.254373 2.203214 2.151474 2.227014 1.997511 1.628392 2.027461
## LAS-Min  2.254025 2.262128 2.144635 2.245310 2.016814 1.641607 2.048312
## LAS-1se  2.309883 2.316721 2.172645 2.287196 1.989073 1.652671 2.074964
## PLS      2.331110 2.285897 2.173540 2.217923 2.038460 1.647188 2.040617
## GAM      2.091092 1.933269 1.882283 2.059659 2.022764 1.631334 2.022538
## Full-Tree 2.868454 2.343721 2.282801 2.712111 2.513663 2.474951 2.741155
## Min-Tree 2.304910 2.393777 2.036022 2.235144 2.075740 1.789475 2.082046
## 1SE-Tree 2.309883 2.130736 1.932816 2.287196 2.056348 1.792475 2.074964
##           8           9          10
## LS       1.892345 2.086754 2.002616
## Step     1.914008 2.063151 2.038700
## Ridge    1.886477 2.069589 1.990547
## LAS-Min  1.892848 2.068131 2.013093
## LAS-1se  1.931517 2.099387 2.044572
## PLS      1.884597 2.060205 2.060292
## GAM      1.694160 1.758533 1.908523
## Full-Tree 1.909727 2.558427 2.378676
## Min-Tree 1.824420 1.938802 2.028302
## 1SE-Tree 1.931517 2.099387 2.044572
```

Average MSPE


```
rowMeans(CV.MSPEs)
```

```
##      LS      Step      Ridge  LAS-Min  LAS-1se      PLS      GAM Full-Tree
## 2.057040 2.081994 2.043605 2.058690 2.087863 2.073983 1.900415 2.478369
## Min-Tree 1SE-Tree
## 2.070864 2.065989
```

Make boxplot


```
boxplot(CV.MSPEs, las = 2, main = "MSPE Boxplot")
```


MSPE Boxplot



```
### Get relative MSPEs and make boxplot
CV.RMSPEs = apply(CV.MSPEs, 1, function(W) W/min(W))
CV.RMSPEs = t(CV.RMSPEs)
boxplot(CV.RMSPEs, las = 2, main = "RMSPE Boxplot")
```

RMSPE Boxplot



```
# Code for best model
```

```
# Initializing environment
```

```
h2o.init(nthreads = -1)
```

```
## Connection successful!
```

```
##
```

```
## R is connected to the H2O cluster:
```

```
## H2O cluster uptime: 2 hours 4 minutes
```

```
## H2O cluster timezone: America/Vancouver
```

```
## H2O data parsing timezone: UTC
```

```
## H2O cluster version: 3.34.0.3
```

```
## H2O cluster version age: 1 month and 21 days
```

```
## H2O cluster name: H2O_started_from_R_dhruv_ygi863
```

```
## H2O cluster total nodes: 1
```

```
## H2O cluster total memory: 0.70 GB
```

```
## H2O cluster total cores: 4
```

```
## H2O cluster allowed cores: 4
```

```
## H2O cluster healthy: TRUE
```

```
## H2O Connection ip: localhost
```

```
## H2O Connection port: 54321
```

```
## H2O Connection proxy: NA
```

```
## H2O Internal Security: FALSE
```

```
## H2O API Extensions: Amazon S3, XGBoost, Algos, AutoML, Core V3, TargetEncoder, Core V4
```

```
## R Version: R version 4.1.1 (2021-08-10)
```

```

y = "y"
x = setdiff(names(data), y)

train.h2o = as.h2o(data)

```

```
## |
```

```

h2o.fit1 <- h2o.gbm(
  x = x,
  y = y,
  training_frame = train.h2o,
  nfolds = 5,
  ntrees = 5000,
  stopping_rounds = 10,
  stopping_tolerance = 0,
  seed = 301471961
)

```

```
## Warning in .h2o.processResponseWarnings(res): early stopping is enabled but neither score_tree_inter
```

```
## |
```

```
h2o.fit1@parameters$ntrees # 41
```

```
## [1] 39
```

```
h2o.rmse(h2o.fit1, xval = TRUE) # 1.278083
```

```
## [1] 1.278757
```

```

split = h2o.splitFrame(train.h2o, ratios = 0.75)
train = split[[1]]
valid = split[[2]]

```

```

h2o.final <- h2o.gbm(
  x = x,
  y = y,
  training_frame = train.h2o,
  nfolds = 5,
  ntrees = 5000,
  learn_rate = 0.01,
  learn_rate_annealing = 1,
  max_depth = 1,
  min_rows = 1,
  sample_rate = 0.75,
  col_sample_rate = 1,
  stopping_rounds = 10,
  stopping_tolerance = 0,
  seed = 301471961
)

```

```
## Warning in .h2o.processResponseWarnings(res): early stopping is enabled but neither score_tree_inter
```

```
##      |
```

```
h2o.final@parameters$ntrees
```


```
## [1] 5000
```

```
h2o.rmse(h2o.final, xval = TRUE)
```

```
## [1] 1.294353
```

```
h2o.varimp_plot(h2o.final, num_of_features = 10)
```

Variable Importance: GBM



```
# Validating model
```

```
test.h2o <- as.h2o(model_test_data)
```

```
##      |
```

```
h2o.performance(model = h2o.final, newdata = test.h2o)
```

##

```
## RMSE: 1.789799
```

```
## RMSLE: 0.1345529
```

```
h2o.predict(h2o.final, newdata = test.h2o)
```

|

```
## 1 12.71176
```

```
## 3 12.84556
```

```
## 5 13.08423
```

##

```
pred.h2o = predict(h2o.final, test.h2o)
```

|

```
# Getting prediction for test_data
```

|

```
## [1] "WARNING: Model metrics cannot be calculated and metric_json is empty due to the absence of the "
```

```
h2o.predict(h2o.final, newdata = test.h2o)
```

|

```
## 1 12.11686
```

```
## 3 13.12212
```

```
## 5 12.61131
```

##

```
## [3000 rows x 1 column]
```

```
predictions = predict(h2o.final, test.h2o)
```

```
## |
```

```
h2o.exportFile(predictions, path = "test.csv")
```

```
## |
```